

Sockets

learn:

- [x] Most and Least significant bit
- [] Endianness
- [] Octets
- [x] two's complement

We can think of sockets as a way to speak to other programs using standard unix file descriptors.

When unix programs do any sort of IO they do it by reading or writing to a file descriptor, which is an integer associated with an open file ("everything in unix is a file!").

Calling `socket()` returns us a socket descriptor, which then allows us to communicate through it using specialized `send` (for sending) and `recv` (for receiving) socket calls.

Stream Sockets

`SOCK_STREAM` are a reliable two-way connected communication streams with guaranteed order. Web browsers use the HTTP protocol, which uses stream sockets to get pages. Uses TCP and IP under the hood.

Datagram Sockets

`SOCK_DGRAM` also uses IP for routing but UDP for transmission. They're called 'connectionless' as we don't need to really keep a connection open.

Structs and Types

Most significant Bit bit position in a binary number having the greatest value. it can also correspond to the sign bit of a signed binary number in one's or two's complement notation, 1 meaning negative and 0 positive. MSB and LSB are indications of the ordering of the sequence of bits in the bytes sent over the wire. MSB means that the MSB will arrive first. LSB means that the LSB will arrive first.

Two's Complement corresponds to a binary signed number representation based on a simple mathematical operation. It is defined as the complement to 2^N (result of subtracting the number from 2^N - taking the one's complement and adding one). This corresponds to the negative version of the original number in a way that both can coexist in a natural way.

`memset(void *str, int c, size_t n)` copies the character `c` (unsigned char) to the first `n` characters of the string (block of memory) pointed to, by the argument `str`.

Endiannes is the ordering or sequencing of bytes of a word of digital data in computer memory storage or during transmission. There are two of them: *big-endian* and *little-endian*. The first stores the most-significant byte of a word at the smallest mem address and the least significant byte at the largest. The other is the inverse. As an example, x86 processors store in little-endian. **In data networking Big-endian is the most common convention**, then giving the name of *network byte order*.

An example of **big-endian** (and so, **network byte order**) is the daily representation of the number 123, with the hundreds (which is the most significant - has the biggest value) at the first position in the memory and the least significant at the last positions.

socket descriptor: int

uint16_t htons(uint16_t): converts host byte order to net byte order

uint16_t ntohs(uint16_t): converts net byte order to host byte order

```
/**
 * Holds socket address information for
 * many types of sockets
 */
struct sockaddr {
    // address family (commonly AF_INET)
    unsigned short sa_family;
    // 14 bytes of protocol address
    // dest address and port number
    char sa_data[14];
}

/**
 * just like sockaddr but specialized
 * for internet stuff. Can also be casted
 * to 'socketaddr' to pass to socket.
 * Note that sin_zero[8] should be set
 * all to zeros via memset to pad the
 * struct correctly to sockaddr.
 */
struct sockaddr_in {
    short int sin_family; // AI
    unsigned short int sin_port; // NBO
```

```

    struct in_addr      sin_addr; // NBO
    unsigned char       sin_zero[8];
}

struct in_addr {
    // 32-bit long, 4 bytes
    // just use the proper conversion
    // function to make it NBO
    unsigned long s_addr;
};

```